

Remotely Operated Vehicle (ROV) and Artificial Intelligence (AI) test results, current status, and future action items

Zachary Randell, Research Scientist

Alex Tanz, Marine Science Interpreter & Diver

Shawn Larson, Curator of Conservation Research

Conservation Programs & Partnerships, Seattle Aquarium

April 22, 2022

Contents

1 Abstract	1
2 ROV testing	1
2.1 <i>September 1st</i>	1
2.2 <i>November 10th</i>	1
2.3 <i>November 12th</i>	2
2.4 <i>December 10th</i>	3
2.5 Next steps re: ROV testing and survey capacity	4
3 Artificial Intelligence (AI) analyses	5
3.1 Percent-cover with CoralNet	5
3.2 Object detection via VIAME	6
3.3 Next steps re: AI analyses	8



Note to the reader: this pdf is from a LaTeX document that contains embedded links, both within (red border, e.g., Fig. 1, for figures and sections) and external (blue border, e.g., url, for relevant websites and videos in Google Drive). When viewing this document offline (e.g., in an adobe reader after downloading the file), if you click on a red box to jump to a figure or section, click and hold *Alt* then *left-arrow* (on PC) or *Command* then *left-arrow* (on Mac) to return to your previous page in the document.

1 Abstract

Our objective is to use the emerging accessibility of Remotely Operated Vehicle (ROV) technology and Artificial Intelligence (AI) software to expand the spatial extent across which we can make inferences about patterns of benthic community structure. When nested within a long-term subtidal monitoring program, we envision such advancements in survey and processing methodology will enhance our ability to identify the processes and factors affecting ecological resilience and kelp-forest persistence through time. In order to provide scientific SCUBA divers an additional benthic survey tool, we are testing and developing ROV methods to survey relatively shallow (5-40m) locations characterized by urchin barrens, and turf, understory, and canopy forming macroalgae. Given the large amount of imagery envisioned, we are training AI algorithms to generate estimates of percent-cover and object (species) abundances at scale.

Regarding current status, we have tested the Seattle Aquarium’s Blue ROV2 on four separate occasions. These tests have largely centered around iterating camera and light placement, and we have developed a custom framework upon which to mount the ROV’s lights. On the AI side, we have established pipelines using CoralNet to analyze percent-cover from photos, and VIAME to detect objects (individuals of conspicuous species) from photos/video. Our immediate next steps involve developing standardized ROV survey protocols and full-scale establishment/testing of AI pipelines. We envision a highly collaborative project, thus this document is intended to set the stage for transparent and reproducible research; open dialogue—your thoughts, feedback, and criticisms—are most welcome.

2 ROV testing

We have tested the ROV on four occasions, with a single *1hr* dive each of the four days. Major points from each of the four days are noted below, and this section closes with an itemized breakdown of future ROV tasks. Finally, big shoutout to Alex Tanz at the Seattle Aquarium for leading the assembly and initial wiring of the ROV, as well as for assisting in all ROV tests that took place at the Seattle Aquarium.

2.1 *September 1st*

We calibrated ROV controls in a holding tank prior to our first flight off the side of the Seattle Aquarium down to 70’ (Fig 1*a*). Controlling the ROV via a Xbox game controller was intuitive and the ROV was very responsive, particularly when “feathering” the controls, i.e., the ROV was responsive to varying strengths of thruster (joystick) input. The QGroundControl software (QGC) that communicates with the ROV was easy to install and can be readily accessed to download logs of ROV activity. No GoPros were used during this first flight.

2.2 *November 10th*

In a first test of the “benthic camera array,” two GoPro HERO 10 cameras were mounted on the payload skid underneath the ROV (Fig 1*b*), and a third GoPro was staged behind the other two and faced forwards. Of the two benthic-facing cameras, one recorded 4K 120FPS (frames per second) video, the other shot 23MP (megapixel) photos every 2 seconds. We once more deployed the ROV off the side of the Seattle Aquarium.

We anticipate sharing the live-feed of the ROV will be an invaluable tool for public outreach and education, particularly for classroom/field trip events with communities along the coast and

in Puget Sound, and also as an additional tool to engage visitors at the Seattle Aquarium. In order to evaluate the ease of sharing the real-time ROV feed, Alex set up a TV, connected it to the ROV laptop, and streamed video from the forward-facing (built-in) camera used to steer the ROV. We also initiated a Zoom meeting from the dedicated ROV laptop, screen-shared, and had Seattle Aquarium personnel in another building open the meeting and view the live ROV video. We are now exploring how to directly connect the ROV video feed as a webcam, i.e., using the ROV camera's IP address. Additional software is required to make this happen, and doing so will provide higher quality video that avoids double-compression in QGC and Zoom. In essence, we want a separate video stream that bypasses QGC and goes straight to Zoom, or whatever other media outlet source is desired, e.g., a local TV.

The benthic camera array worked well in terms of *obtaining* video and time-lapse photos. However, as you can see in Fig 2a—c, a shadow is prominent in the middle of the frame, and the left-third of the frame is not well illuminated. These lighting effects were not a surprise given that the light configuration on Nov 10th was the “default” configuration, which maximizes light illumination directly in front of (and not below) the ROV (Fig. 1b). Despite suboptimal lighting, the latest-generation GoPros demonstrated a high-degree of adaptability to light variability, and the image-stabilization performed extremely well. You can see this in Vid 1., a short video from the benthic-facing camera. This is a 4K video, though your media player will downscale it, particularly when viewing the clip in Google Drive (though you can download the original 4K file).

2.3 November 12th

Following November 10th's test and the presence of shadows/suboptimal illumination, I changed light placement such that the lower of the two lights were attached to the payload skid (a relatively small adjustment). I traveled to Whidbey Island where Ken Collins was kind enough to host and take me and the ROV out on his vessel. We launched out of Freeland County Park, anchored at the northeast corner of Holmes Harbor ($48^{\circ} 5' 6.2016''$ N, $122^{\circ} 31' 18.4002''$ W), and flew the ROV above an eelgrass bed, and along soft sediment down to 85' (Fig. 2d—f, Vid. 2).

Regarding lighting, while I succeeded at providing more uniform lighting across the camera's field of view, I inadvertently also succeeded at creating a large amount of backscatter by illuminating the water column underneath the cameras. For example, note how the image with eelgrass appears out of focus (Fig. 2f) relative to photos from Nov 10th (Fig. 2a—c). Furthermore, in Vid. 2) note the faint, approximately biradial shadow present (most visible in the short clips along soft sediment); this stemmed from the light's beam hitting the frame upon which the lights were mounted.

Others lessons from Nov 12th: the ROV performed well in a moderate current. In particular, the software governing ROV movement compensates incredibly well for the varying forces acting upon ROV position at any one point in time. BUT, despite being able to point the ROV towards a heading (via the compass on our laptop display), current still had a significant effect upon realized ROV position, e.g., flying the ROV north with a current running west does not put the ROV where one expects it to be. This first field deployment away from the relative shelter of the Seattle Aquarium illustrated the importance of incorporating an acoustic GPS tracking system into our ROV operations. Such a tracking system will provide explicit knowledge of—and control over—ROV position, crucial for basic operations and the execution of repeatable (i.e., consistent survey location) and standardized (consistent area covered) benthic surveys.

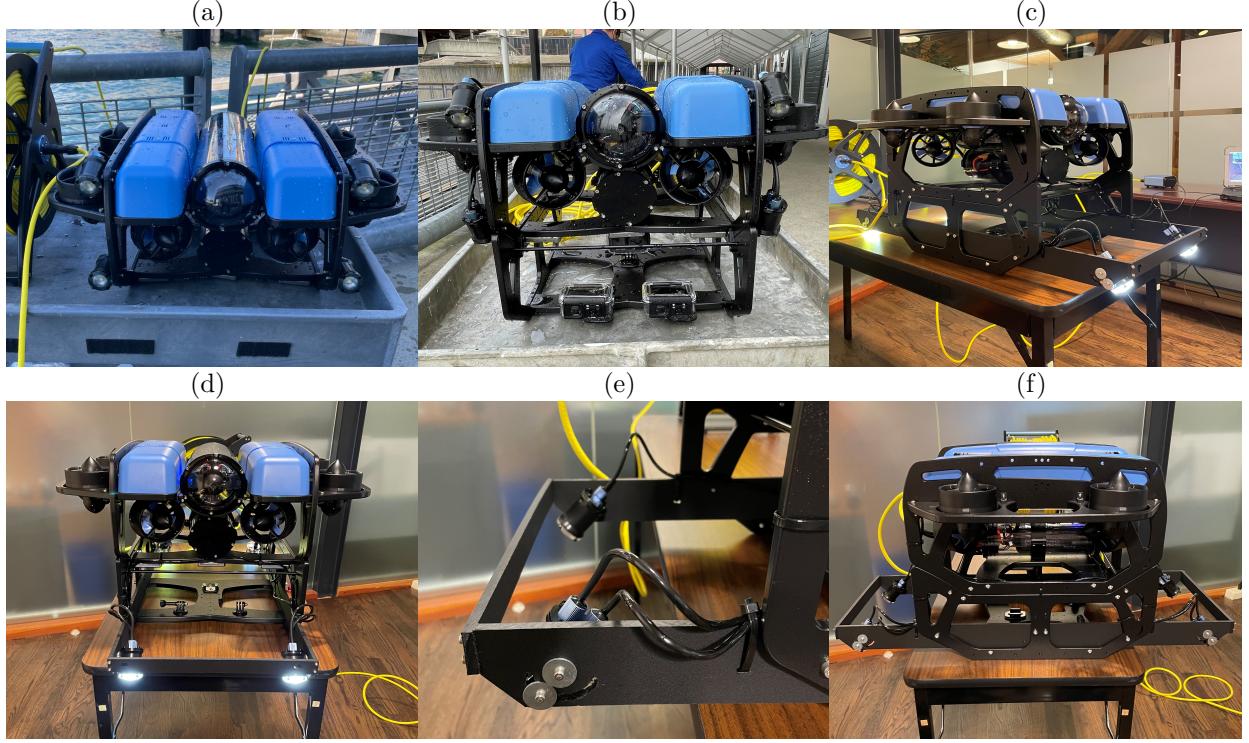


Figure 1: (a) The ROV as it was flown Sep 1st following assembly. (b) Nov 10th we added the payload skid underneath the ROV, along with two GoPro HERO 10 cameras facing downwards. (c) The current version of the ROV includes HDPE panels constructed to provide a framework maximizing the distance between the ROV lights and the GoPro cameras. (d) The new light framework does not increase the lateral profile (i.e., entanglement risk) of the ROV. (e) Pitch angle of the lights can be adjusted on the new frame. (f) Sideways view illustrating the increased fore and aft dimensions with the custom frame.

2.4 December 10th

Following the November tests, it was clear that the pre-drilled options for mounting the four lights were not suited to properly illuminate the benthos underneath the ROV. I therefore obtained HDPE panels—the same material comprising the frame of the ROV—to create a custom structure upon which to mount the lights. This custom framework allowed us to stage the lights fore and aft around the ROV (Fig. 1c). Our framework did not change the width of the ROV (Fig. 1d), thus minimizing any additional entanglement risk. To ensure the frame did not “grab” stipitate kelp between the two new panels, we closed off the fore and aft framework with a small HDPE bar affixed perpendicular to the primary HDPE panels (Fig. 1e). Light angle can once more be adjusted (Fig. 1e). Finally, in order to further minimize backscatter, the cameras were moved from the front of the payload skid to the middle, maximizing the distance between the camera and lights.

To test this frame Alex and I briefly flew the ROV off the side of the Seattle Aquarium on Dec 2nd. No effects of the frame upon ROV maneuvering or buoyancy were observed. As you can see in Vid. 3, the placement of the lights along the new frame indeed minimizes backscatter. I was a wary about fully powering the lights (as doing so had made the backscatter worse on Nov 12th), but as we had successfully addressed the causal source of the backscatter, in hindsight our imagery from Dec 1st would have benefited from increased light power—this is because the outward angle of the

lights reduced how much light reached the seafloor beneath the ROV. That being said, as noted above, we do not see backscatter, and the lighting is more uniform across the frame. Note however the black bar that forms when the ROV gets close to the substrate. This is actually “good” in the sense that the spacing of the lights—in addition to the angling of the lights—minimized light illumination of the water column beneath the cameras. When the ROV descends close enough to the seafloor that the fore and aft light beams do not intersect, the black bar is the result. We simply need to adjust the angle of the lights to match our desired ROV altitude. In short, this framework appears to have addressed the fundamental lighting issues encountered previously, and further testing will allow us to optimize light settings.

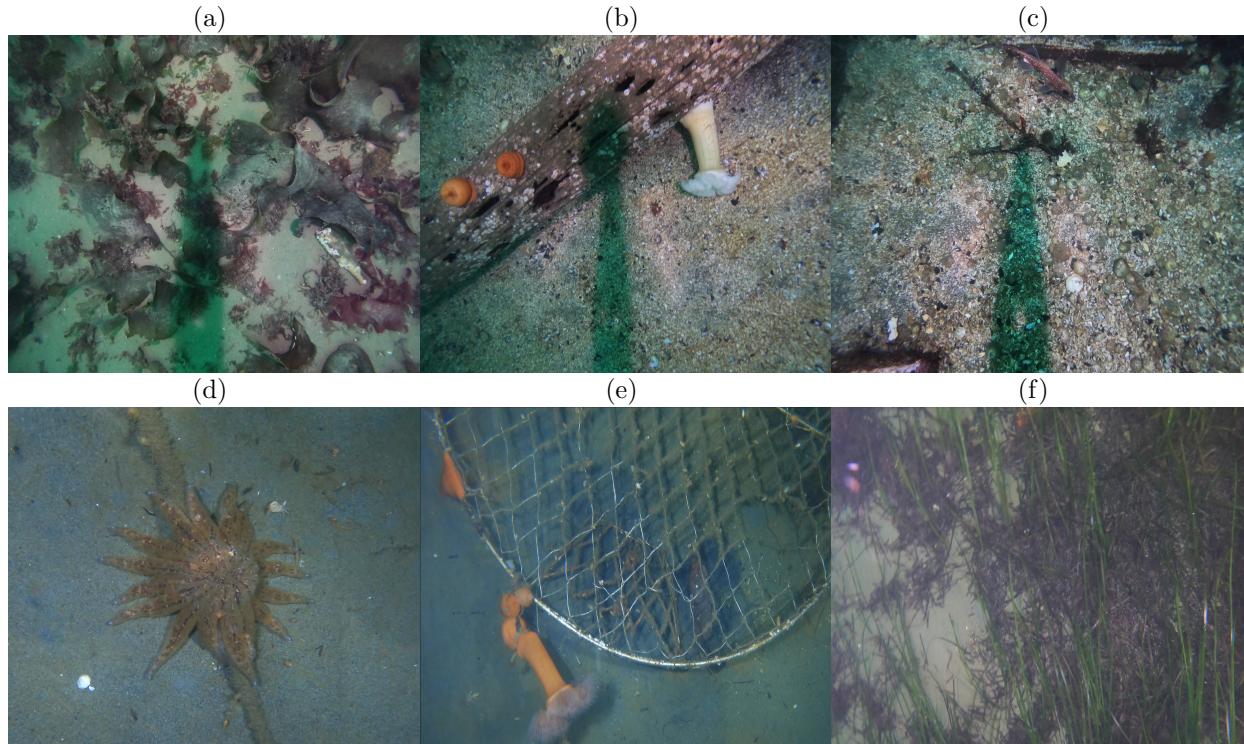


Figure 2: (a—c) 23MP photo from 2sec time-lapse from Nov 10th. Note the shadow and poorly illuminated left third of the frame due to the placement of the lights relative to the frame of the ROV. Despite the suboptimal lighting, the resolution is clear enough for, e.g., percent-coverage image analysis of red algae, sugar kelp, green algae, sandy substrate in (a), or object detection of anemones and ratfish, in (b) and (c), respectively. (d—f) Screenshots of video offshore of Whidbey Island from Nov 12th. (e) An abandoned crab pot containing crabs and pycnopodia was observed at a depth of 85’—no sign of this pot is visible on the surface. (f) Eelgrass imagery demonstrating the potential for percent-coverage analyses to detect: live eelgrass, dead eelgrass, and soft sediment categories. Note how unfocused (f) is relative to (a—c); this is due to backscatter from changing the position of the lights relative to the camera.

2.5 Next steps re: ROV testing and survey capacity

- Cover fore and aft “corners” with small pieces of HDPE paneling to prevent light from washing out the forward-facing camera used to steer the ROV.
- Add ability to adjust lights laterally (in addition to forward and reverse pitch, which can

currently be adjusted).

- Ensure lights can be locked into place to eliminate the possibility of inadvertently nudging their placement.
- If desired, fabricate a framework using aluminum.
- Incorporate and test Ping Sonar Altimeter that will provide real-time feed of ROV altitude relative to seafloor. This will enable us to obtain imagery at a consistent scale through precise control of ROV altitude.
- Incorporate and test WaterLinked’s Underwater GPS G2, an acoustic GPS tracking system. This will provide a real-time feed of the ROV’s position that will allow precise navigational control, as well as the ability to repeatably survey consistently sized transects.
- Within the Dome exhibit at the Seattle Aquarium, fly the ROV alongside a diver in real-time to test different light angles, light power, and ROV altitudes; precisely identify an optimal configuration for image acquisition.
- Possibly within the Dome but also in the field, lay out a meter tape and test varying ROV speeds in order to identify optimal survey speed. Similarly, test how camera settings and ROV altitude affect the precise width of imagery attained.
- Use width and speed information to calculate (and verify) survey area per unit effort, e.g., m^2 imaged / 1hr ROV operation.

3 Artificial Intelligence (AI) analyses

We have trained preliminary algorithms to (1) classify percent-coverage from photos using CoralNet and (2) identify objects (conspicuous and readily identifiable individuals) from photos/video using VIAME. As the results from VIAME are new and have yet to be presented/discussed among our collaborators, I focus upon VIAME here and largely skim over CoralNet. My intent is not to comprehensively delve into the details of software or the constituent models/algorithms, but rather to provide a brief overview of VIAME and the results of a preliminary proof-of-concept analysis of time-lapsed imagery from an urchin barren.

3.1 Percent-cover with CoralNet

To briefly recap our progress with CoralNet, we are using the online portal to analyze still images and obtain measurements of percent-cover. Our preliminary proof-of-concept pipeline was principally focused on simple categories such as urchin vs rock vs water vs diver (gross simplifications, but these categories sufficed for proof-of-concept). CoralNet is intuitive and user friendly, and our initial results were very encouraging (Fig. 3). Big shoutout to Ken Collins for first tackling CoralNet in the latter parts of 2020 and early 2021—these preliminary analyses were instrumental to first orientating ourselves to AI image analysis. Upon acquisition of suitable ROV-derived benthic imagery, we will train CoralNet to calculate percent-coverage of, e.g., red algae, understory brown algae, sessile and colonial/aggregate invertebrates such as sponges, tunicates, hydroids, and bryozoans, along with substrate type such as soft-sediment, mud stone, shell debris, cobble, and hard substrate.

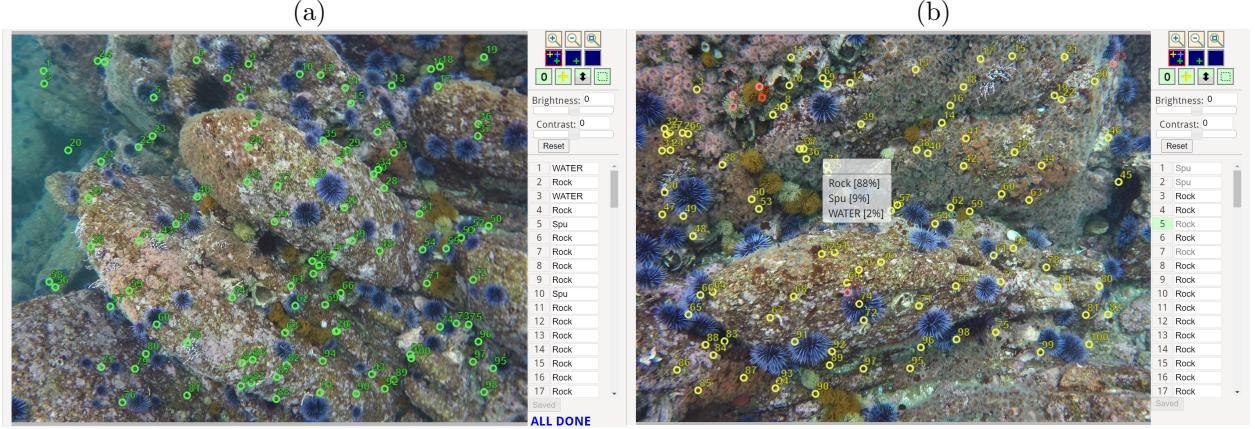


Figure 3: (a) Screenshot of an image that has been manually annotated, i.e., an observer has selected a cover category for each of the 100 randomly placed green circles. (b) Screenshot of an image that has been annotated by the trained algorithm, i.e., the computer makes its best guess for what the category underlying each circle is. Yellow circles have been confirmed by an observer, and the red point that is selected (upper left) displays the percent confidence of the algorithm’s predictions. Selecting a point also provides the opportunity to correct the prediction, which in turn increases the accuracy of future predictions.

3.2 Object detection via VIAME

Video and Image Analytics for Multiple Environments (VIAME) is an open-source AI toolkit developed by Kitware with funding from NOAA. Here are links to an overview document, a website linking to most all VIAME information, and the VIAME online interface. VIAME is open-source on GitHub. Thus far I have opted to use the desktop version of VIAME, opposed to the online interface. This is because the desktop version provides access to all project files, whereas the online version automates the process and renders project files inaccessible. While there are extra considerations (and responsibilities) when managing project files, I opted for this so that we have the option to make our algorithms open-source, e.g., on GitHub. This will ensure anyone can download (but not change) our files, i.e., anyone can make use of the trained algorithm.

Shoutout to Olivia Rhoades, a postdoc with the Hakai Initiative in British Columbia, and Kathleen MacGregor, Aquatic Science Technician with Fisheries and Oceans Canada, Quebec, for providing me video and time-lapse photography (e.g., Fig. 4a,b) of urchins in an urchin barren. These urchin-barren video and time-lapse photos were suited for investigation with VIAME, and they allowed me to develop the following proof-of-concept pipeline using VIAME.

I first manually annotated half of a sequence of time-lapse images (Fig. 4c). This “training” phase provided the information used to compile an algorithm—given our annotations—that can then be applied to other, unseen data. Annotations in VIAME involve creating a box around each observation of an object (e.g., an urchin, a crab, etc.), and correcting the position of the box as the object moves through the time-lapse or video. VIAME has an option to automatically track the object in question, so following an individual through an image sequence is quick and easy. See Vid. 4 for a screen-recording that scrolls through a sequence of manually-annotated images. (And note VIAME works perfectly well for “single-detections”, i.e., a single observation of an object). To test VIAME’s ability to differentiate conspicuously visible versus cryptic individuals, I broke the category “urchin” into *urchin* (the urchin is clearly visible) and *covered_urchin* (urchins with decorations, e.g., rocks, algae, etc., on their test/spines). Many *covered_urchins* had no actual

spines or parts of their test visible, and as they often looked identical to small rock piles, they were ideally suited to test VIAME's ability to differentiate cryptic species. I furthermore annotated the movements of two very small sea stars to test how well VIAME handled tiny species comprised of relatively few pixels (Vid. 5).

With annotations in place, I compiled the annotations into a trained algorithm (a pipeline). I opted to train a SVM (support-vector machine) model. One of the advantages of supervised SVM models is their ability to compile rather quickly—it took 17:03min for my annotations to be compiled into a usable algorithm. In contrast, Deep Learning or Convolutional Neural Network (CNN) models could take several days if not a week to initially train from a set of annotations. (VIAME can train CNN models, and we will explore these once we have a more robust set of training annotations). Applying our trained SVM algorithm to unseen data (the other half of the time-lapse sequence) took 11:25min. Given that we only had 100 individuals annotated across approximately 80 frames (photos), the algorithm performed surprisingly well when applied to unseen imagery (Fig. 4d). See Vid. 6 for a screen-recording of the classified output, i.e., of output from applying the algorithm to unseen data; I scroll through frames from the time-lapse, and also scroll through % confidence, a setting that filters predictions based on how confident the algorithm is regarding each detection.

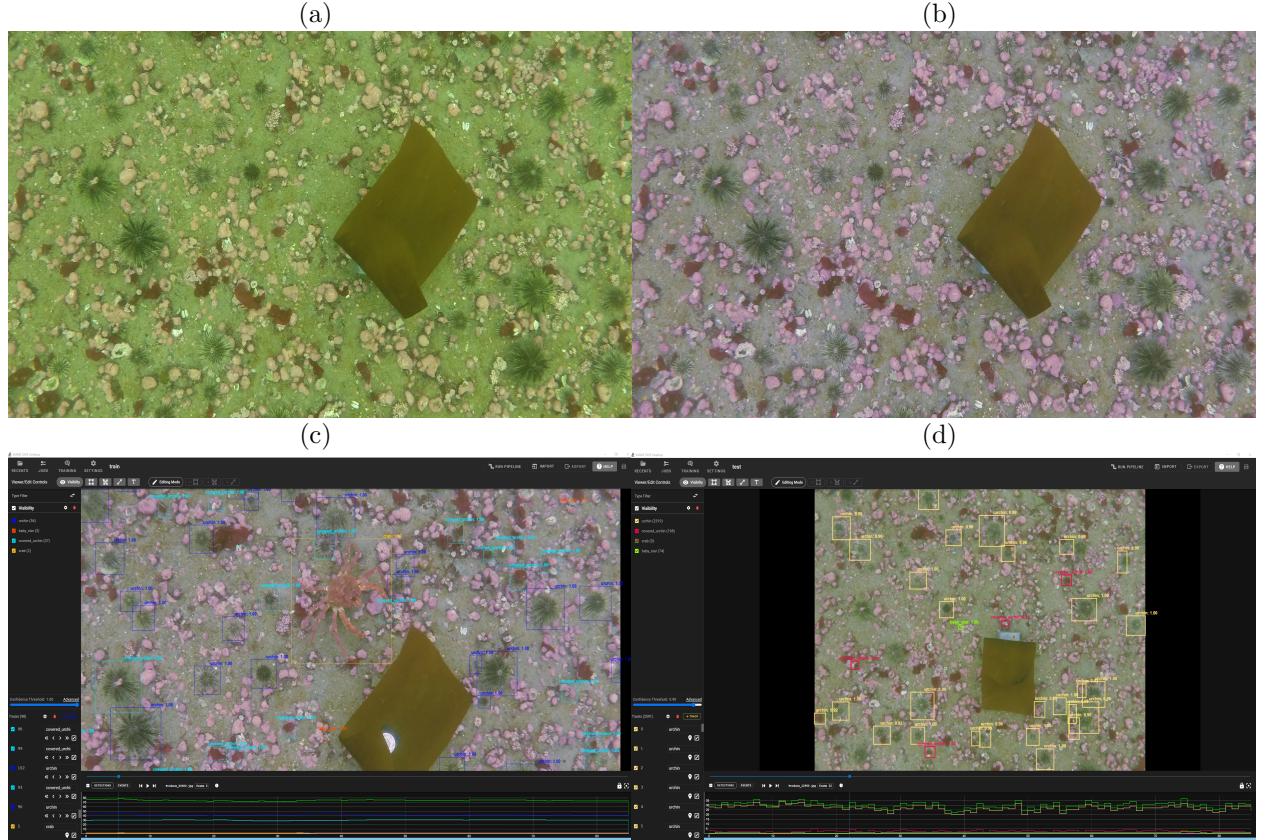


Figure 4: (a) Screenshot of raw image as I received it from Kathleen MacGregor, who staged a Go-Pro that took a photo every 1min. (b) Screenshot of same image after adjusting the white balance. (c) Screenshot of VIAME while creating manual annotations, i.e., “informing” the computer what an urchin is, what a crab is, etc. (d) Screenshot of the trained algorithm applied to unseen data, i.e., the algorithm’s attempt at identifying species given the manual annotations we provided.

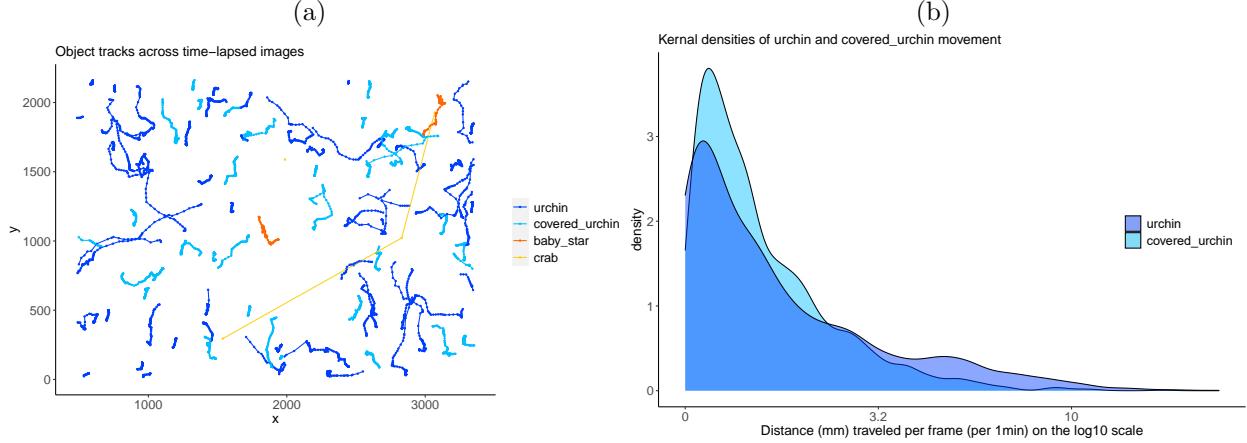


Figure 5: (a) Movement tracks of individuals from the four categories of species: *urchin*, *covered_urchin*, *crab*, and *baby_star*. (Note a large *crab* was observed in three sequential images before moving out of frame, hence the two long and straight lines connecting each of the three observations). (b) Movement patterns summarized with kernel densities for *urchin* and *covered_urchin*, with the former exhibiting high frequencies of longer movements (with step-length in *mm* and log base-10 transformed).

Finally, data from annotating and then applying an algorithm are easy to work with via exported CSV files. As an example, Fig. 5 presents the tracking data generated from annotating the urchin-barren imagery. Each “line” represents a single individual (urchin, crab, etc.) tracked across the entirety of the training sequence (Fig. 5a). Fig. 5b visualizes the kernel densities of individual step-lengths for *urchins* and *covered_urchins* (with step-length distance in *mm* that have been log base-10 transformed). We see *covered_urchin* exhibit a higher frequency of short steps, whereas *urchins* predominantly exhibit the longer step-lengths. If we wanted to test for differences between these kernel densities (distributions), we could use a two-way Kolmogorov-Smirnov Test. Finally, note that we could also incorporate urchin size information (based on the size of VIAME’s bounding box around each observation) to model the urchin-size vs step-length relationship.

3.3 Next steps re: AI analyses

- In order to establish version-control and enable accessible research, I have created a GitHub repository (a repo) with all necessary files to run the VIAME algorithm trained here. This GitHub repo—https://github.com/zhrandell/AI_proofOfConcept—is currently set to “public”, so anyone can access and download the files. However, these files are only usable if one has downloaded VIAME. Furthermore, this repo will be more useful when I can upload imagery so that one can reproduce the process of training an algorithm (I have not yet received permission to publicly release the images I used). Additional future action items include creating tutorial content: with a Zoom recording, I will talk through and screenshare the steps of “pulling” from and “pushing” to a GitHub repository, downloading VIAME, training an algorithm, applying a trained pipeline, etc.
- Next steps with CoralNet include using ROV-obtained images to create and classify real categories of interest, e.g., red algae, brown algae, colonial/aggregate invertebrates, and substrate type (as described in the CoralNet subsection 3).
- Next steps with VIAME (similar to CoralNet) involve training an algorithm on ROV-derived

imagery. Additionally, there are other photogrammetry tools in VIAME (such as the ability to generate photo mosaics) that will be worth exploring.

- Finally, once a robust training set of images have been annotated in VIAME, we will compile deep learning or CNN models.